

# Notes for presentation

Here are some suggested notes for the presentation. Feel free to make changes and additions but remember that they expect all group members to talk for about 5 minutes, give or take.

## **Structure:**

### *Problem definition*

- Introduction of project, problem and the Gym environment (Clement)

### *Solution*

- Artificial Neural Networks and how we use them in our project (Andreas)
- Genetic algorithms. What are they and why use them? (Kristian)
- GA – Crossover Algorithm (Ronny)
- GA – Mutation Algorithm (Kristian)
- Short description of code / implementation? (Andreas?)

### *Results*

- Results and demonstration (Ivica)

The requirements for the presentation is that it “**cover[s] the problem definition, solution and results**”. And so, each section should not be too focused on theory, but be a combination of how the algorithms work in theory, and how we have implemented / used them in this project.

# Introduction of project, problem and the Gym environment (Clement)

## Keywords

- Using artificial neural network to interact with a simulated environment
- Using genetic algorithms to improve ANNs over multiple generation
- Very short description of OpenAI GYM
- Description of the CartPole environment.
  - Problem to be solved

## Suggestion for text

In this project we tackled the problem of training artificial neural networks to perform a task in a simulated environment.

The training happens by using multiple generations of many individual networks. Genetic algorithms are used to develop favorable traits, and the best resulting networks are the outcome of the process. Hopefully, these will be able to perform the task on a satisfactory or high level.

OpenAI Gym is a toolkit for training reinforcement learning algorithms and there are multiple environments and tasks to test an agent's learning ability against. The one we have made use of in this project is called the CartPole environment. This is a 2D game where a cart move along a track. The cart can move to the left or the right on the screen. On top of the cart a pole is placed in a hinge and balanced. The goal of the game is to keep the pole upright by pushing the cart towards the right or towards the left, to stop the pole from falling.

[The slide should contain

a picture/video of CartPole environment

and an illustration of observation array]

The game or environment is at any point described with 4 observations:

- The location of the cart
- The velocity of the cart
- The angle of the pole
- The velocity of the pole at the tip

There is no need to know about what has happened previousy, the agent only needs to know the current state of the enviornment, described by these four numbers.

The artificial neural networks take these four observations as input and then decides on the next action to take, which is to apply a unit force in the left or right direction. The game ends either when the pole falls below a certain angle, when the cart hits the boundaries of the game, or when the maximum number of timesteps have passed. Every timestep the network receives an observation

and it must decide on an action. It is not possible to do nothing; a force must be applied in either direction.

Initially, the environment starts at a value close to zero. There is a small random number given to all the four observations to get the game started. We also assume that the first action performed is a random action, so a push either in the left or the right direction, determined by a randomizer function. From there on, it is our agent's task to move the cart to balance the pole.

Originally, the game was meant to stop after simulating 500 timesteps, but to make sure the agents become very good at balancing, we let the simulations run much longer. The limit we set was 10 000 timesteps.

The simulation of this game happens to every artificial neural network in every generation until the desired number of generations have passed. We can choose whether to render the environment, meaning display what happens in every step. We will not do this while training the networks, as this will take too much time, but at the end of the presentation, we will perform a running through the program and display the best network's performance.

# Artificial Neural Networks and how we use them in our project (Andreas)

## Keywords

- Description of how ANNs work
  - Weights, biases, activation function
- Implementation of ANNs in this project
  - Input nodes
  - Hidden layer
  - Output
  - Specific implementation
    - Class
    - MLPClassifier
    - ReLU
    - Etc.
- Descriptions of generational improvement of ANNs
  - Genetic algorithms explained in detail by others

## Suggestion for text

# Genetic/Evolutionary Algorithms and the Mutation Algorithm (Kristian)

The current plan is that you introduce genetic algorithms and talk about why we might want to use it in general and for this project. Then, after Ronny talks about the crossover algorithm, you describe the mutation algorithm.

We divided it like this because the crossover algorithm takes longer to explain, and these two parts should take up about the right amount of time combined.

## Keywords

- Short description of theory of genetic algorithms
  - Type of evolutionary algorithm
  - Models the algorithm as a biological entity, with genes
- Why genetic algorithms?
  - Advantages
  - Disadvantages

- How we use genetic algorithms for this project
  - Individual weights and biases are genes
  - Sets of weights and sets of biases are chromosomes
  - An ANN is seen as an individual
  - New ANNs are made by selecting best ANNs to parent new, child ANNs
    - Probability based on reward
    - Crossover algorithms used to create children ANNs
  - Mutation algorithm used to introduce new information

**Suggestion for text**

## **GA - The Crossover Algorithm (Ronny)**

The current plan is that you talk about the crossover algorithm after Kristian introduces genetic algorithms, but before he describes the mutation algorithm.

### **Keywords**

- Uses the weights and biases from the parent ANNs
  - Generates child ANNs
    - New weight and bias configuration
- Our implementations
  - Single-point crossover
  - Two-point crossover
  - Uniform
  - Unravel matrix or keep shape?

**Suggestion for text**

## **Short description of code / implementation? (Andreas?)**

We did not decide on this in our discussion, but I thought I (or maybe someone else, if their part is too short) could give a quick description of the code and how it runs. This would literally be 1 minute or less.

### **Keywords**

- Quick description of classes
- Quick description of running the code
  - Initializing generation of ANNs
  - Running simulation for each ANN
    - Reward
  - Generating new generation of ANN
    - Using crossover and mutation
  -

### **Suggestion for text**

## **Results and demonstration (Ivica)**

You can decide what you want to say and show for the results like we discussed, but we should probably cover:

- Short description of why we chose the different methods / implementations in the algorithms – if you don't do any testing for one of the choices, maybe say a word or two about why you prefer one over the others
- Why we chose to partial\_train using env.reset() and avoided partial\_training for every step
  - Nice if you say a sentence on why you think env.observation\_space.sample() gives bad results, like you explained at the meeting
- Results
  - How many generations until solution, on average, or something like that
  - Other metrics of your choosing
  - Remember to mention that even a good ANN can still fail
- Comparing best ANN, average ANN – seems like a demand in the assignment text?
- Live demonstration running code